

Tutorial No. 4

jupyter Untitled4 Last Checkpoint: 2 minutes ago

File Edit View Run Kernel Settings Help

📁 + ✂ 📄 📌 ▶ ■ ↺ ⏭ Code ▼

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc_curve, auc

[2]: diabetes = load_diabetes()
X, y = diabetes.data, diabetes.target
y_binary = (y > np.median(y)).astype(int)

[3]: X_train, X_test, y_train, y_test = train_test_split(X, y_binary, test_size=0.2, random_state=42)

[4]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

[5]: model = LogisticRegression()
model.fit(X_train, y_train)

[5]: ▾ LogisticRegression
LogisticRegression()

[6]: y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: {:.2f}%".format(accuracy * 100))

Accuracy: 73.03%
```

```
[6]: y_pred = model.predict(X_test)
      accuracy = accuracy_score(y_test, y_pred)
      print("Accuracy: {:.2f}%".format(accuracy * 100))
```

Accuracy: 73.03%

```
[7]: print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
      print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Confusion Matrix:


```
[[36 13]
 [11 29]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.77	0.73	0.75	49
1	0.69	0.72	0.71	40
accuracy			0.73	89
macro avg	0.73	0.73	0.73	89
weighted avg	0.73	0.73	0.73	89

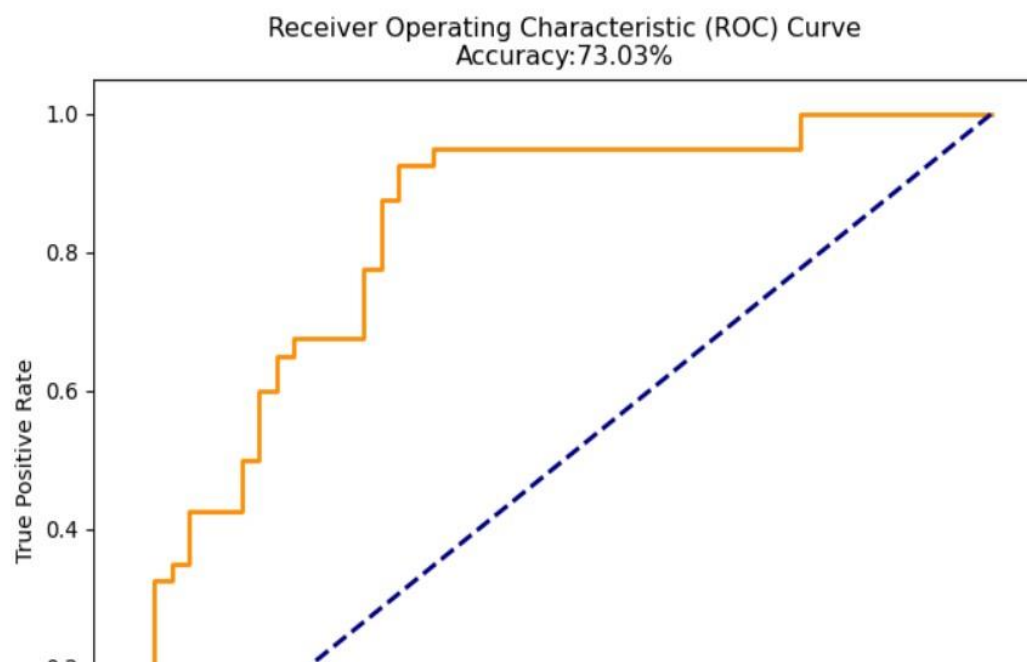
```
[8]: plt.figure(figsize=(8, 6))
      sns.scatterplot(x=X_test[:, 2], y=X_test[:, 8], hue=y_test, palette={0:'blue', 1:'red'}, marker='o')
      plt.xlabel("BMI")
      plt.ylabel("Age")
      plt.title("Logistic Regression Decision Boundary\nAccuracy: {:.2f}%".format(accuracy * 100))
      plt.legend(title="Diabetes", loc="upper right")
      plt.show()
```

```
plt.ylabel("Age")
plt.title("Logistic Regression Decision Boundary\nAccuracy: {:.2f}%".format(accuracy * 100))
plt.legend(title="Diabetes", loc="upper right")
plt.show()
```

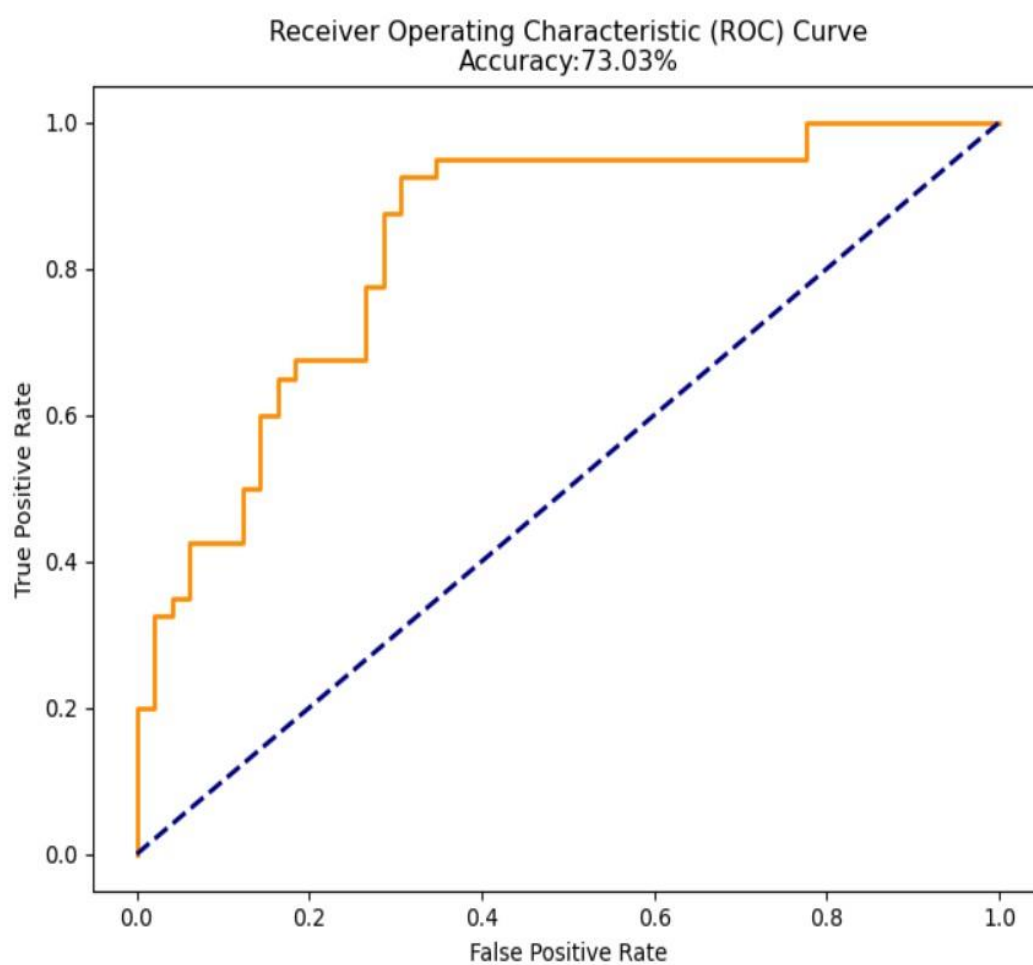


```
[9]: y_prob = model.predict_proba(X_test)[: , 1]
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='Random')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve\nAccuracy:{:.2f}%'.format(accuracy * 100))
```

[9]: Text(0.5, 1.0, 'Receiver Operating Characteristic (ROC) Curve\nAccuracy:73.03%')



```
[9]: Text(0.5, 1.0, 'Receiver Operating Characteristic (ROC) Curve\nAccuracy:73.03%')
```



[]: